

## Deep Learning Basics

GO DATA DRIVEN



### Goal

- Become familiar with common components of a neural network.
- Understand the effects that different hyperparameter choices can have on the model.

## Deep Learning Basics

GO DATA DRIVEN



### Program

### Architecture:

- Fully-connected layers
- Activation function
- Bias
- Batch normalization layer
- Dropout layer

### Hyperparameters:

- loss function
- gradient descent type
- weight updates
- weight initialisation

## Architecture


We will discuss the following neural network components

- Fully-connected layers
- Activation function
- Bias
- Batch normalization layer
- Dropout layer

## Fully-Connected layer

All inputs of one layer connected to every activation unit of the next layer.


Also known as *Linear* or *Dense* layer



## Activation

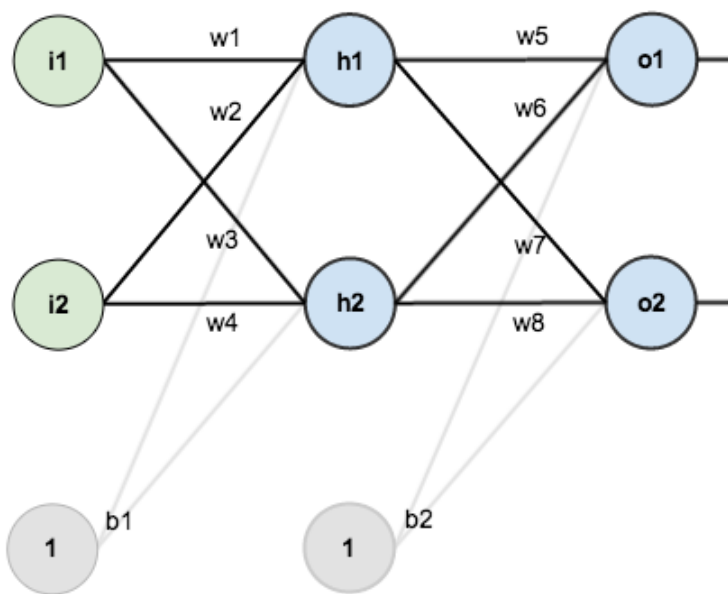
Introduces non-linearity into the network.

No trainable parameters.



## Bias

An additional parameter that allows you to shift the input to the activation function to the left or right (which may be critical for successful learning).



## Batch Normalization

Batch normalization ([Lofte et al.](#))

- Normalize the layer inputs with batch normalization.
- This helps to ensure all layers activated in near optimal “regime” of the activation functions.
- Since the gradients’ dependency on the scale of the weights is reduced, it allows us to use higher learning rates,
- which means training is accelerated, as less iterations are required to converge to a given loss value.

<b>Input:</b> Values of $x$ over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$ ;	
Parameters to be learned: $\gamma, \beta$	
<b>Output:</b> $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$	
$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$	// mini-batch mean
$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$	// mini-batch variance
$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$	// normalize
$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$	// scale and shift

## Batch Norm (notes)

Batch Norm learns 4 parameters


- $\beta$
- $\gamma$
- running mean  $\mu$  (for inference stage)
- running variance  $\sigma^2$  (for inference stage)

## Normalization (further reading)

- Weight normalization ([Salimans et al.](#))
- Layer normalization ([Ba et al.](#))


## Regularization

The great flexibility of neural networks makes them very powerful, however this comes at the price of easily overfitting of the data.




## Dropout

- "Drop" neurons in the network with probability  $p$  (every mini-batch/epoch)
- No trainable parameters



(a) Fully-Connected Neural Network



(b) Applying dropout with  $p \simeq .5$

## Dropout

- Computing the gradient is done with respect to the error, but also with respect to what all other units are doing. Therefore certain neurons may fix the mistakes of other neurons.
- Dropout prevents over-reliance on a subset of the neurons in a layer
- every neuron becomes more robust

## Hyperparameters

We shall discuss the follow hyperparamter choices,

- loss function
- gradient descent type
- weight updates
- weight initialisation

## Loss function

**Classification:**

- Binary cross entropy:  $L_{binary} = \frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$
- Categorical cross entropy:  $L_{categorical} = \frac{1}{N} \sum_{i=1}^N \sum_{j=i}^c [y_{ij} \log(\hat{y}_{ij})]$

**Regression**

- Root mean squared error (RMSE):  $L_{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^n (y - \hat{y})^2}$


## Loss function

- **Multi-class classification**
  - **softmax** output layer with **categorical** cross-entropy and **one-hot** targets.
- **Binary or multi-label classification**
  - **sigmoid** output layer with **binary** cross-entropy and **binary** vector targets.
- **Regression**
  - **linear** output layer with **RMSE**
  - Not performing? Try **discretizing output** through binning. Otherwise, go for a different learning algorithm.

Problem Type	Output Type	Final Activation Function	Loss Function
Regression	Numerical value	Linear	Mean Squared Error (MSE)
Classification	Binary outcome	Sigmoid	Binary Cross Entropy
Classification	Single label, multiple classes	Softmax	Cross Entropy
Classification	Multiple labels, multiple classes	Sigmoid	Binary Cross Entropy


## Gradient descent

- Stochastic gradient descent - feed a single datapoint in at each pass.
- Batch gradient descent - feed in the whole batch of data at each pass.
- Mini-batch gradient descent - feed in a group of data at each pass.



## Gradient descent


We update our model using gradient descent.



$$w' = w - \eta \frac{\partial J(w)}{\partial w}$$

## Gradient descent

## Computational resources per pass




## Passes required to find good $w, b$ values


- **Forward pass:** a data sample is passed forward through the network to determine a prediction
- **Backward pass:** recursively compute the error backwards from the last layer following the chain-rule and update the weights w.r.t. the known target output.
- **Epoch:** training the neural network with all the training data for one cycle.

## Weight updates

**Learning rate:** small value  $\eta$  typically between 1.0 and  $10^{-6}$



Large learning rate: Overshooting.




Small learning rate: Many iterations until convergence and trapping in local minima.

## Weight updates





**Learning rate:** small value  $\eta$  typically between 1.0 and  $10^{-6}$



## Weight updates

**Momentum:** take into account the gradient estimation of the previous batches

*SGD with momentum, Nesterov momentum*



## Momentum (further reading)

The main difference is in classical momentum you first correct your velocity and then make a big step according to that velocity (and then repeat), but in Nesterov momentum you first making a step into velocity direction and then make a correction to a velocity vector based on new location (then repeat).

i.e. without momentum:

$$vW(t+1) = - \text{scaling} * \text{gradient\_F}( W(t) )$$

$$W(t+1) = W(t) + vW(t+1)$$

Classical momentum:


$$vW(t+1) = \text{momentum} * vW(t) - \text{scaling} * \text{gradient\_F}(W(t))$$

$$W(t+1) = W(t) + vW(t+1)$$

While Nesterov momentum is this:

$$vW(t+1) = \text{momentum} * vW(t) - \text{scaling} * \text{gradient\_F}(W(t) + \text{momentum} * vW(t))$$

$$W(t+1) = W(t) + vW(t+1)$$



[source.](#) [source](#)


## Weight updates

**Adaptive learning rate:** adapt the learning rate based on the gradient history (removing the dependency on hyperparameter choice).

*AdaGrad, AdaDelta, RMSprop*

**Momentum & adaptive learning rate:** *Adam, Nadam*

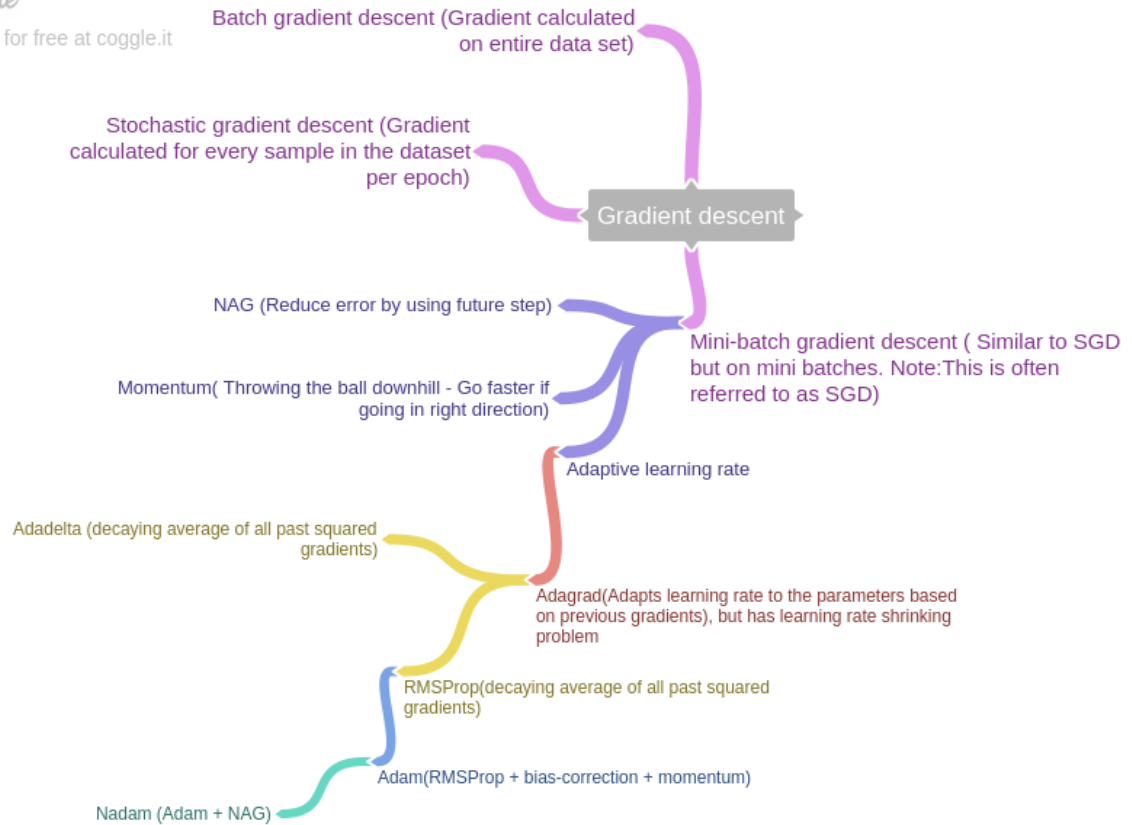
More on this later!



## Optimizers (further reading)

[UvA notes on optimizers](#)


## Weight updates



## Weight initialization

There are a few contradictory requirements:


- Weights need to be small enough magnitude → Otherwise output values explode
- Weights need to be large enough magnitude → Otherwise signal too weak to propagate



## Weight initialization

**Naive approaches:** All zero


Every hidden unit will get zero signal. No matter what the input was, the output would be the same!



## Weight initialization

**Naive approaches:** All constant (e.g. all 1.0)

- Input to each neuron in a layer will be the same,
- therefore the update each neuron in a layer receives will be the same,
- this will prevent different neurons in a layer from learning different things.



## Weight initialization

**Solution:** Break symmetry with a random initialization.

- Xavier or Glorot init:  $w \sim \sqrt{\frac{2}{n_{in} + n_{out}}} \cdot N(0, 1)$  (Glorot et al.)
- He init:  $w \sim \sqrt{\frac{2}{n_{in}}} \cdot N(0, 1)$  (He et al.)

where  $n_{in}$  and  $n_{out}$  represent the number of in and outgoing connections

## Recap Hyperparameters

- **Gradient descent:** dependent on data, usually mini-batch
- **Error function:** dependent on the type of problem; influences final activation
- **Weight updates:** dependent on data, usually Adam is a good optimizer choice
- **Weight initialisation:** He or Xavier

## Hyperparameters Questions

- Which is computationally more expensive: SGD, batch gradient descent or mini-batch gradient descent?

batch gradient descent

- Why is cross-entropy preferred over e.g. classification error (/accuracy)?

Accuracy is not a continuously differentiable function of the weights!

- Which optimizer combines both an adaptive learning rate with momentum?
- Why would you not initialize the weights to 0?
- In what cases would you use He/Xavier initialization over random initialization?

## Summary

In this notebook we covered,

### Architecture:

- Fully-connected layers
- Activation function
- Bias
- Batch normalization layer
- Dropout layer

### Hyperparameters:

- loss function
- gradient descent type
- weight updates
- weight initialisation



# GO DATA DRIVEN

## Neural Networks in Keras

Let's put our knowledge to practice.

### Keras basics

# GO DATA DRIVEN